AN ETHICAL AND EMOTIONALLY INTELLIGENT SOCIAL MEDIA PLATFORM USING RESPONSIBLE AND EXPLAINABLE AI

Project Id: 25-26J-137

Project Proposal Report

Ullandupitiya U.P.L.I.

IT22318152

B.Sc. (Hons) Degree in Information Technology Specialized in Data Science

Department of Computer Science and Software Engineering

Sri Lanka Institute of Information Technology

Sri Lanka

August 2025

AN ETHICAL AND EMOTIONALLY INTELLIGENT SOCIAL MEDIA PLATFORM USING RESPONSIBLE AND EXPLAINABLE AI

Project Id: 25-26J-137

Project Proposal Report

Ullandupitiya U.P.L.I.

IT22318152

Supervisor: Mr. Samadhi Rathnayake

B.Sc. (Hons) Degree in Information Technology Specialized in Data Science

Department of Computer Science and Software Engineering

Sri Lanka Institute of Information Technology

Sri Lanka

August 2025

DECLERATION

We declare that this is our own work, and this proposal does not incorporate without acknowledgement any material previously submitted for a degree or diploma in any other university or Institute of higher learning and to the best of our knowledge and belief it does not contain any material previously published or written by another person except where the acknowledgement is made in the text.

Name	Student ID	Signature
Ullandupitiya U.P.L.I.	IT22318152	

The above candidates are carrying out research for the undergraduate Dissertation under my supervision.

Signature of the supervisor: Date: 28/8/25

(Mr. Samadhi Rathnayake)

ABSTRACT

Today, global communication has completely transformed by social media platforms, enabling billions of people to engage, share ideas, and get connected. Along with all the positives, some significant challenges like misinformation, harmful content, cyberbullying, and fake or bot like accounts have also raised with this rapid growth. These inauthentic profiles are often automated or operated with deceptive intentions which may spread false information, harmful content, malicious narratives, facilitate scams, and weaken the user trust. Manual moderation or content-based filtering are the traditional approaches to detect fake or bot like accounts which are reactive, easily evaded, and resource intensive, making them ineffective to meddle with large scale, evolving threats.

This research centers on Behavior-Based Fake Account Detection module within an integrated, emotional intelligent, and ethical social media platform that leverages responsible and explainable AI. While the overall system address features such as cyberbullying detection, emotion-driven reactions with ethical search filtering, and transparent and explainable friend recommendations, this module specifically focuses on fake account detection while collaborating with other modules to create a unified harm mitigation strategy.

This approach shifts from traditional content-based analysis to behavioral and network-based indicators such as friend request trends, posting frequencies, profile completeness, account lifespan, interacting patterns, and graph metrics like clustering coefficient. Unsupervised models including DBSCAN and Isolation Forest are used to detect anomalies without relying on labeled data. Appropriate moderation actions are enforced accordingly with the implemented dynamic risk scoring mechanism based on its detected behavioral anomalies, while offering a clear justification for the applied moderation actions using explainable AI techniques.

The anticipated result is a proactive, scalable, and explainable detection framework that leverages integrity of the platform, lessens the manual moderation, and uphold user trust ensuring both ethical and transparent governance.

Keywords: Fake Account Detection, Behavior-based Analysis, Unsupervised Learning, Network Metrics, Explainable AI

TABLE OF CONTENTS

DECLERATION i

ABSTRACT ii

LIST OF FIGURES v

LIST OF TABLES vi

LIST OF ABBREVIATIONS vii

1 INTRODUCTION 1

1.1 Background & Literature Survey 3

1.2 Research Gap 6

1.3 Research Problem 7

2 OBJECTIVES 10

2.1 Main Objective 10

2.2 Sub-Objectives 10

3 METHODOLOGY 12

3.1 Data Collection and Preprocessing 12

3.2 Behavior-Based Fake account Detection Module 13

3.3 Component Architecture 14

3.4 Evaluation Metrics 15

3.5 Tools and Technologies 16

3.6 Work Breakdown Structure 16

3.7 Gantt Chart 17

4 PROJECT REQUIREMENTS 18

4.1 Functional Requirements 18

4.2 Non-Functional Requirements 18

4.3 Expected Test Cases 19

5 Commercialization and Budgeting 21

REFERENCES 22

LIST OF FIGURES

Figure 3.1 - Component Architecture Diagram..... 15

Figure 3.2 - Work Breakdown Structure of Behavior-Based Fake Account Detection Component 16

Figure 3.3 - Gantt Chart 17

LIST OF TABLES

Table 1.1 - Research Gap Table..... 7

LIST OF ABBREVIATIONS

Abbreviation	Description
AI	Artificial Intelligence
API	Application Programming Interface
AUC-ROC	Area Under the Curve – Receiver Operating Characteristic
LSTM	Long Short-Term Memory
MAU	Monthly Active Users
REST	Representational State Transfer
US	Unites States
XAI	Explainable Artificial Intelligence

1 INTRODUCTION

In the present world, social media has become an integral part of the modern society, linking billions of people, enabling to engage, share ideas, and get connected together. As of July 2025, the number of social media users around the world has reported as **5.41 billion** which represents **65.7 percent** of the global population. The study shows that a typical social media user engages with about **6.84** different social media platforms monthly and spends around **18 hours and 46 hours** in social media per week on average [1]. These platforms have transformed, revolutionized, and reshaped the personal relationships, business operations, and news distribution. Hence new avenues for civic engagement were provided, establishing a cultural, economic, and political influence surpassing the history of digital communication.

Despite its transformative impact with many benefits, the rapid growth of social media has also brought significant challenges that risks both individual well-being and the stability of the communities. Spread of misinformation, circulation of harmful and abusive content, manipulation of algorithmic recommendation, and rise of fake or bot like accounts are the among key concerns. Such inauthentic profiles distort publish disclosure, promote malicious narratives, and weaken user trust whether they are automated, semi-automated or human controlled under false identities. Meta's Community Standards Enforcement Report states that the estimated percentage of fake accounts was approximately 3 percent of worldwide MAU on the platform Facebook during the 1st Quarter of 2025. Also, it shows that they have taken actions on more than 1 billion of inauthentic accounts within the same quarter of 2025 [2]. These illustrations underscore the vast scale of the problem as well as the difficulty of addressing it effectively.

The anticipations of harmful online behavior have gone far beyond the digital space, leading to severe phycological distress, and even loss of life in extreme situations. The research shows that, compared with nonvictims, those who have experienced were cybervictimization were almost twice as likely to self-harm, exhibit suicidal behaviors, attempt suicides, or have suicidal thoughts [3]. As of 2020, and according to the Center

for Disease Control and Prevention, suicide was the third leading cause of death among adolescents in the US, from 1999 to 2006 with the online harassment incidents highlighted as a key contributing factor in most of the cases [4].

UNICEF Global Survey shows that, approximately one in three young people across 30 different countries have experienced some form of online bullying, and closely **20 percent** of children skip school as these affect the mental health of children and young people turn to substances like drugs, alcohol or violent behavior to deal with their phycological and physical pain [5] . In Sri Lanka, **over 3000** complaints have reported regarding online harassment, fraudulent social media profile incidents, illustrating the increase of reported cases [6].

In addition to harassment, fraud, scams, and thefts are frequently driven by fake or inauthentic user accounts. The recent statistics showed by Crime in England and Wales reveal that **3.9 million** fraud incidents were reported in the year ending September 2024, with a notable **19 percent** increasement showing how these online threats are evolving and how widespread it is [7]. The research done by Stanford Internet Observatory underscores how fake profile networks spread false information causing a persistent threat across social media platforms [8].

To address these challenges, this research proposes an Ethical and Emotional Intelligent Social Media Platform designed with responsible and explainable AI (XAI), maintaining principles of safety, transparency, fairness, and user phycological well-being. The platform combines four AI integrated modules:

1. **Cyberbullying Detection and Behavioral Profiling -** Detect harmful or abusive content without relying on manual reporting.
2. **Emotion-Aware Reaction System and Ethical Search Filtering** – Customize reactions with the emotional tone and filter search queries which are unsafe, unethical.
3. **Explainable Friend Recommendation Engine** – Reliable connections are suggested with clear explanation for each suggestion.

4. **Behavior-Based Fake Account Detection** – The focus of this study, detect fake or bot like accounts and enforce appropriate moderations accordingly without depending on manual interventions.

To combat online threats ongoing, the behavior-based Fake Account Detection module fights against inauthentic profiles by focusing on behavioral patterns and indicators that harder to fake such as friend request patterns, posting frequency, engagement patterns, account lifespan, and profile completeness. With the anomaly detection models like Isolation Forest, DBSCAN, it enables to detect anomalies without the need for labeled datasets assembling the dynamic risk scoring mechanism to determine the appropriate moderation actions which may ranging from account monitoring to suspension. With responsible AI-driven explanations ensuring the transparency, and fairness, while output of this module will enhance the resilience and user trust within the platform.

In summary, this research focus on creating and delivering a proactive, scalable, and explainable detection module that protect integrity of the platform as well as mitigates the real-world harm. Behavior analysis, moderation pipeline, explainable AI are integrated into an ethical architecture which will propose a transparent, and a trustworthy solution ensuing a safer online environment.

1.1 Background & Literature Survey

Social media platforms allow billions of people to connect globally revolutionizing human interactions, its openness also resulted for many opportunities for misuse through the rise of fake and bot-like accounts. Those inauthentic accounts which may fully automated, semi-automated, or managed in human under false identities are frequently used to manipulate public opinion, spread false information, commit scams and fraud, and promote abusive behavior. Meta reported that they have taken actions more **1.5 billion** fake accounts in the third quarter of 2022 and **99.6%** of the accounts were detected automatically rather than getting flagged through user reports. [9]

In the early approaches, content-based analysis was the primary approach relied on to detect fake or bot-like accounts, where suspicious and inappropriate texts, posts or images were counted as inauthentic activities. Although these methods can identify spams clearly, such approaches can be easily evaded through paraphrasing, image manipulation, and generative texts. Ferrara et al. (2016) describes this challenge as an "arms race" between defenders and attackers [10]. To surpass these challenges, researchers have been shifted more towards behavioral and network-based indicators such as friend request velocity, account interactivity patterns, account lifespan, posting frequency, and network metrics like clustering coefficients. Furthermore, it allows platform to identify evolving threats by longitudinal tracking of user behavior including slow infiltration and compromised accounts which one-off checks may not sufficient or fail to capture.

Most of the current behavioral detection modules are remain on supervised models heavily, which requires large and up-to-date labeled training data. Such data becomes outdated quickly, may be vulnerable to issues like bias, and expensive. Many of these systems act only as advisory tools, leaving the enforcement of final moderations under manual control which can be inconsistent and inefficient at large scale. In contrast, an unsupervised, explainable AI-based framework offers greater promise by analyzing the evolving user behavioral patterns, enforcing appropriate moderation measures automatically, and provisioning transparent explanations enhancing scalability, resilience, and user trust.

Over the past few years, academic and industry researchers have investigated to come up with different detection frameworks by investigating multiple approaches to surpass the challenges, ranging from content-driven solutions to hybrid and unsupervised behavioral solutions. Collectively, these studies provide valuable insights into the strengths and limitations of existing studies, offering a foundation to develop more advanced, transparent, and effective detection system.

The UnDBot, a novel that propose an unsupervised, interpretable framework for detecting social bots. UnDBot has built upon structural information theory which

capture various aspects of social bot behaviors such as Posting Type Distribution, Posting Influence, and Follow-to-Follow Ratio. These behavior/profile signals are utilized to construct multi-rational graph, aiming to model the social user behavior relevance and to discover the long-distance correlations between users. They have introduced a novel approach to optimize heterogenous structural entropy to identify bot communities without labels, claiming its effective performance across multiple datasets. Key strengths of this framework include independence from labels, interpretable graph-level logic reasoning, and flexibility across platforms since its features are not tied to a single dataset. [11]

Zouzou and Varol's in 2024 research investigates the suspicious following patterns among all the followers of a social media account to combat the challenge of fake-follower campaigns by examining the irregularities in how followers get connected with certain accounts. With the detection of unsupervised learning, it captures the synchronized surges of follower activity that inflate the credibility artificially. The strengths that showcase in this approach is its ability to generalize without needing the labelled data and captures large-scale collective strategies which may remain hidden at the individual account-level. This study highlights that focusing on coordinated dynamics is important rather than treating suspicious accounts individually. [12]

Ianus, a framework built by WeChat and introduced to detect the fake account at its registration stage. Here, unusual device usage, repeated IP addresses, and large number of new accounts are the focused key system checks to identify any suspicious activity patterns. The key strength of this approach is that it has been tested in the real world, handling millions of sign-ups daily. The powerful strategy of stopping fake accounts before they misuse the platform ensures a protected environment showing that early intervention can be highly effective on reducing future threats. [13]

MulBot, an unsupervised framework which counts on user actions as time-series data. Usual or normal behavior is learnt by LSTM autoencoder and then the groups of accounts with bot-like rhythms are identified by clustering. The key advantage in this approach is that this framework allows to detect bots that change slowly over time

focusing on time-behavior patterns. This research highlights how the combination of deep learning and time-series analysis techniques can uncover advanced automation strategies that simple, static methods often miss. [14]

Explainable Fake Account Detection with SHAP (2024) states that this study navigates a crucial turning point with the help of Explainable AI technique (SHAP) to justify the results of an XGBoost model, providing an explained, transparent detected results. It focuses on behavioral signals including user posting frequency and follower ratios, giving strong performance while also explaining globally and locally for each classification. The strength in this approach lies on its transparency, allowing moderators to clearly understand which feature influenced what decision. This powerful strategy enhances an ethical, transparent system making it more trustworthy for human oversight. [15]

1.2 Research Gap

Although the reviewed literature contributes significantly to fake or bot-like account detection, each address only a fragment of the overall problem. UnDBot and Zouzou & Varol shows the importance and its strength of unsupervised detection, but it lacks continuous monitoring or dynamic risk scoring at the individual level. Strong real-world performance is showed by Ianus, but it only focuses on the register time detections and leaving the later behavior unchecked. MulBot framework captures evolving user activity well, yet a justification for the results or automatic enforcement is not discussed. The SHAP-based study improves transparency, but it follows supervised learning approaches which remains relying on labeled datasets and does not integrate dynamic risk scoring or automated moderations.

Table 1.1 - Research Gap Table

	UnDBot [11]	Zouzou and Varol [12]	Ianus [13]	MulBot [14]	Explainable Fake Account Detection [15]	**Proposed System**
Behavioral & profile (non-content)	✓	✓	✓	✓	✓	✓
Unsupervised / weakly-supervised	✓	✓	✗	✓	✗	✓
Temporal monitoring	✗	✗	✗	✓	✗	✓
Dynamic risk scoring	✗	✗	✗	✗	✗	✓
Explainable AI	✗	✗	✗	✗	✓	✓
Automated moderation	✗	✗	✗	✗	✗	✓
Real-time integration	✗	✗	✓	✗	✗	✓

Taken together, these studies reveals that there is no existing solution combines unsupervised detection, ongoing monitoring, dynamic risk scoring, explain ability and automated moderation into a single framework. The proposed system addresses this gap offering an end-to-end component that detect suspicious behaviors, assigns dynamic risk scores, enforces appropriate moderations automatically, and provides clear justifications for the decisions in real-time.

1.3 Research Problem

Social media platforms still struggle with ongoing challenge of inauthentic accounts like fake, and bot-like accounts as well as groups of coordinated actors. These accounts

are used to spread false information, manipulate public opinion, run scams, commit fraud, and weaken overall trust eroding user confidence. Although platforms have come up with many safeguard approaches, attackers adapt, making it difficult to capture by the detection systems. The main challenge is separating real users from advanced inauthentic accounts that closely reflects the usual human behavior.

Most of the current approaches follow content faltering or manual oversight. Content-based systems can be easily evaded through paraphrasing, manipulating images, or generative text, while manual oversight is slow, inconsistent, and impractical at large scale. More advanced approaches follow behavioral and network signals like posting frequencies, friend-request rates as indicators to detect anomalies, but many approaches remain supervised, relying on labeled data which is not the best option as it is costly, biased, and quickly outdated. Furthermore, existing approaches only perform detection and stop at it, leaving enforcement under human,
which leads delays and inconsistent results.

A further limitation is that most detection systems treat accounts as static, going through single checks rather than assessing them over time. Fake accounts evolve gradually and only later shows suspicious or abnormal patterns. In addition, most of the solutions do not provide explainable justifications, leaving enforcement actions appear arbitrary and weakening user trust in the platform's governance.

The research problem can therefore define as, existing fake account detection approaches are fragmented and reactive, lacking a unified framework that brings together unsupervised learning approach, continuous tracking, dynamic risk scoring, explain ability., and automated enforcement. This gap leaves social media platforms vulnerable to emerging threats and weakens fairness and transparency in their moderation.

To address this, the proposed component delivers a comprehensive approach that monitor behavioral patterns and profile metadata over time. It has the mechanism of dynamic risk scoring which can adapt as behavior changes, and proportionate

modifications are enforced automatically with risk score levels. Importantly, it incorporates explainable AI strategy to make each, and every decision is transparent and ethically explained.

In summary, the problem is that existing detection systems fall short because they do not deliver a complete, adaptive, and transparent defense. These shortcomings are directly addressed by the proposed system, offering an end-to-end framework that enables continuous monitoring, automated enforcement, and strengthen user trust in moderation processes.

2 OBJECTIVES

2.1 Main Objective

The main objective of this study is to develop an explainable, automated AI-driven framework which also follows unsupervised learning to detect and manage fake or bot-like accounts on social media platforms. This framework performs continuous monitoring over time, dynamic risk scoring, and automated enforcements for moderate actions. Integrated explainable AI techniques provides ethical justifications for each decision, ensuring accountability and fostering user trust. The aim of this objective is to build a complete, scalable solution that addresses the challenges of existing detection systems, strengthening platforms against evolving threats from inauthentic accounts.

2.2 Sub-Objectives

1. **Behavioral Data Analysis:** To collect, clean, and preprocess data regarding user activity such as friend request patterns, posting frequency, user interacting patterns, and profile completeness, to organize a structured feature set that can be effectively applied to anomaly detection models.

2. **Model Development:** To build, train, and test multiple unsupervised anomaly detection models such as Isolation Forest and DBSCAN and evaluate models' performances in identifying suspicious accounts without relying on labeled data.

3. **Dynamic Risk Scoring:** To establish a framework that assigns and continuously updates the risk score values along with the evolving user behavior, ensuring that abnormal behaviors are identified promptly as they appear over time.

4. **Automated Moderation Pipeline:** To implement an automated pipeline that connects detection outcomes straightly to appropriate moderation actions such as

feature restrictions, shadow bans, or permanent suspension, minimizing the manual intervention in moderation.

5. **Explainable AI Integration:** To embed explainable AI strategies to ensure a transparent, fair, and ethically accountable platform by delivering interpretable explanations for the enforced moderation actions.

6. **System Integration and Deployment:** To connect the detection and moderation component with platform's backend and frontend, enabling detection of suspicious accounts in real-time, automatic enforcement, and administrative oversight.

3 METHODOLOGY

This methodology outlines the structured approach applied to design, implement, and evaluate the Behavior-Based Fake Account Detection Module. The framework is organized to detect and manage fake or bot-like accounts with its objectives without depending on user complaints or manual labels. Unsupervised anomaly detection, dynamic risk scoring, automated moderations, and explainable AI techniques are integrated together emphasizing the scalability, fairness, and transparency. A modular architecture also enables seamless integration into the MERN (MongoDB, Express, React, Node.js) based broader social media platform, making the system practical, and deployable in real-world settings.

3.1 Data Collection and Preprocessing

Data collection and preprocessing provide a solid foundation for the behavior-based fake account detection module. With the limited and biased nature of labeled data, this study leverages a combination of simulated behavioral data and publicly accessible user activity datasets to model both genuine and suspicious account behaviors. The main focus is on extracting interaction patterns and profile characteristics that capture how users typically behave within a social media environment.

Key behavioral features considered include:

- **Friend request activity** – count of requests sent, acceptance ratios, and frequency patterns.
- **Profile completeness** – level of profile completeness, including profile picture, bio, and other additional metadata fields.
- **Posting frequency** – average number of posts, likes, and comments over specific time intervals (daily/weekly).
- **Engagement patterns** – user interactions through likes, shares, and comments.

To prepare the dataset for analysis, several preprocessing steps are carried out:

- **Normalization** – Standardizing values (ex: request volumes, posting rates) to ensure comparability across accounts.
- **Feature Engineering** – Aggregating activities into time windows (ex: hourly, daily, weekly) to capture evolving behavior patterns.
- **Noice Reduction** – Eliminating outliers or data inconsistencies that has introduced by erroneous or incomplete data.
- **Transformation** – Converting raw behavioral indicators into a structured numerical feature set suitable for unsupervised anomaly detection modeling.

Through these steps, a clean and a structured dataset will be produced by the system that accurately represents user behavior, offering a reliable foundation for detecting anomalies.

3.2 Behavior-Based Fake account Detection Module

The design of the Behavior-Based Fake Account Detection is focused here which applies unsupervised machine learning techniques to detect user behavioral anomalies.

Model Selection

This module adopts unsupervised learning approach surpassing the limitations of labeled fake account datasets. The best detection method is chosen after considering and comparing different algorithms. Genuine users are grouped together by clustering-based approaches, whereas isolation-based approaches focus on capturing rare or extreme behaviors. The final model is selected based on its detection reliability, adaptability to emerging patterns, and feasibility for real-time use.

Dynamic Risk Scoring

A risk score is updated continuously and assigned for every account. The score increases when any abnormal or suspicious activity is observed (ex: sending excessive friend requests with low acceptance rate) and also score decreases when the behavior

stabilizes. Thresholds are defined what, and when the moderation actions should be triggered.

Automated Moderation Pipeline

The moderation actions are triggered corresponding to its assigned risk score.

- **Low Risk:** Normal behavior, no actions are triggered.
- **Medium Risk:** Temporary restrictions, feature limitations may enforce such as friend request limitations, posting frequency limitations.
- **High Risk:** Severe enforcements, such as shadow banning, permanent account suspension.

This automated pipeline reduces reliance on manual moderation ensuring enforcements timely and proportionately.

Explainability

User trust and transparency is promoted by the explainable AI (XAI) strategies integrated in the module strengthening user confidence within the platform. These provides clear justifications for moderation outcomes by highlighting its reasons.

Ex: The system may indicate that an unusual high posting frequency along with an incomplete profile metadata led to the enforcement decision.

3.3 Component Architecture

The component is structured as a layered pipeline:

Input Layer → Processing Layer → Decision Layer → Output Layer

- **Input Layer:** Raw behavioral logs are collected and normalized. (Friend requests, activity counts, profile metadata)
- **Processing Layer:** Feature extraction is performed, and anomaly detection approaches are applied to capture suspicious behavior. (Isolation Forest, DBSCAN)

- **Decision Layer:** Dynamic risk scores are generated and assigned, evaluated them against thresholds, and integrates explainable AI strategies to justify enforced moderations.
- **Output Layer:** Actionable moderation responses are delivered. (Warnings, restrictions, shadow bans, and account suspensions)

Component Architecture Diagram

The component architecture diagram depicts the end-to-end pipeline, starting with user interactions and ending with automated moderation, with explainable feedback loops ensuring transparency.

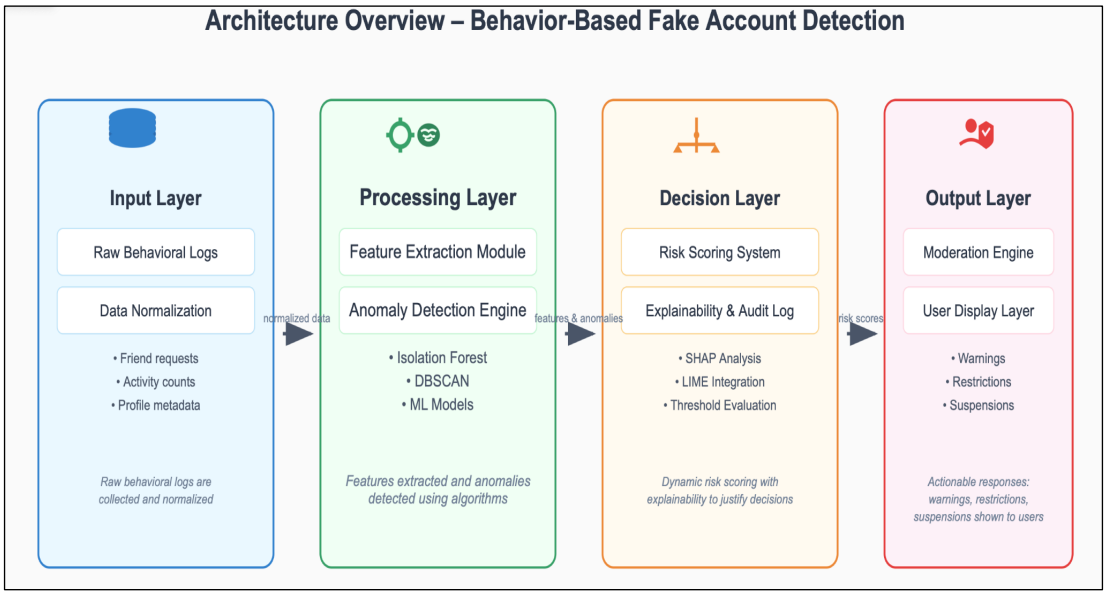

Figure 3.1 - Component Architecture Diagram

3.4 Evaluation Metrics

The performance of the trained anomaly detection models and moderation pipeline is assessed using:

- **Precision & Recall:** To evaluate the detection accuracy and false positives.
- **F1-Score:** Balances the precision and recall measures.
- **AUC-ROC Curve:** To assess the separability between normal and suspicious users.

- **Detection Consistency:** Model performance is compared across different time windows.
- **User Trust Feedback:** Assess transparency and acceptance of justifications for moderation actions.

3.5 Tools and Technologies

- **Development Stack:** MERN (MongoDB, Express, React, Node.js).
- **Machine Learning:** Python (Scikit-learn, PyCaret), Jupiter Notebook
- **XAI Libraries:** SHAP, LIME.
- **System Integration:** REST APIs and WebSocket streams for real-time detection.

3.6 Work Breakdown Structure

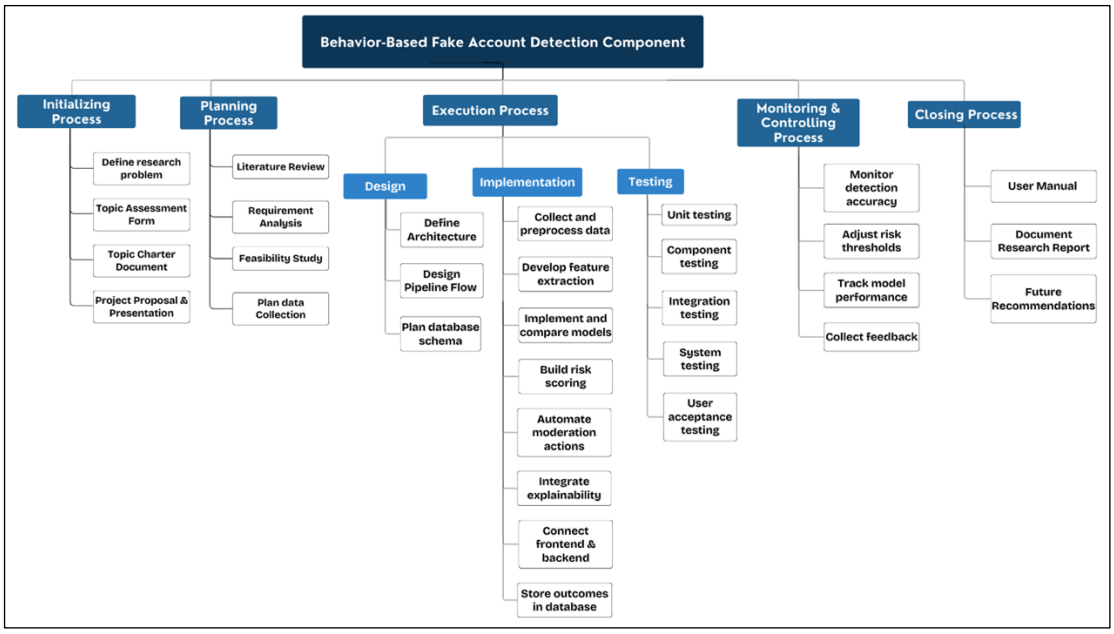

Figure 3.2 - Work Breakdown Structure of Behavior-Based Fake Account Detection Component

3.7 Gantt Chart

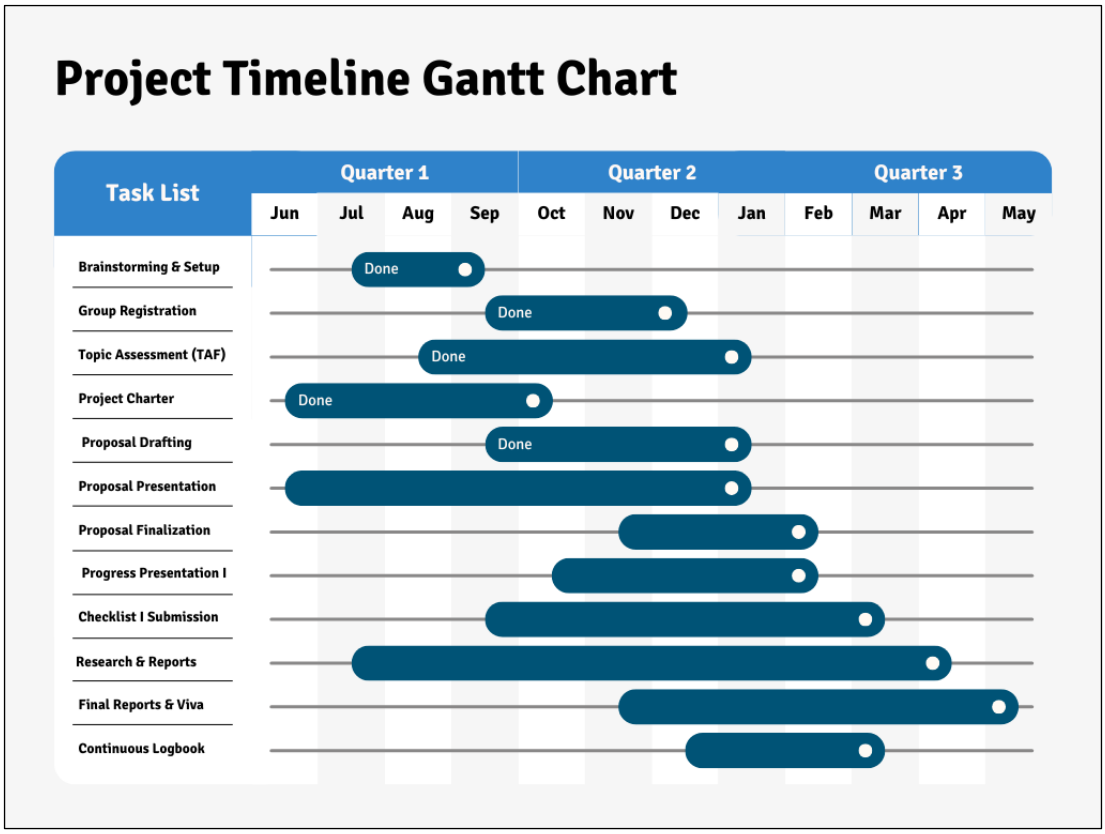

Figure 3.3 - Gantt Chart

4 PROJECT REQUIREMENTS

4.1 Functional Requirements

1. The system should collect and preprocess behavioral and profile signals such as friend request activity, posting habits, interaction patterns, account lifespan, and profile completeness into a structured feature set.
2. The system should follow unsupervised learning to detect anomalies using models such as Isolation Forest, DBSCAN, and assign a dynamic risk score to each account.
3. The system should monitor user behavior continuously and update risk scores with evolving user activity over time.
4. The system should enforce moderations (warn, restrict, shadow ban, ban) automatically based on configurable thresholds.
5. The system should provide human-interpretable justifications for all automated enforcement actions.

4.2 Non-Functional Requirements

1. **Performance** – The system should process behavioral inputs and assign risk scores in near real-time, so that suspicious activity is captured and addressed without noticeable delay.

2. **Scalability** – The system should have the capability to handle the increasing number of users and high volume of events without degradation of its accuracy and speed.

3. **Reliability** – The system should not fail or get stuck throughout the process, ensuring detection and enforcement process remain dependable under all circumstances with minimal downtime.

4. **Security & Privacy** – The system should ensure the protection of sensitive user data through restricting access, encrypting, and securing storage, while ensuring only the behavioral features essential for analysis.

5. **Availability** – The system should stay accessible and functional at all the times, ensuring uninterrupted tracking and enforcement for every active account.

4.3 Expected Test Cases

Testing will be carried out to ensure that behavior-based fake account detection framework functions correctly and satisfies the defined requirements. Both model accuracy and system performance are assessed through different testing phases

Testing Accuracy of the Models

- Dataset for training – 70% of the data from the dataset for model training.
- Dataset for validation – 15% of the data from the dataset for model validation.
- Dataset of test – 15% of the data from the dataset for model testing.
- Model performance and detection reliability will be measured by applying appropriate evaluation metrics.

Testing the Program in Accordance with Testing Phases

1. **Unit Testing**

Each module (data preprocessing, feature extraction, anomaly detection, and risk scoring) will be tested separately and confirm each module has no errors.

2. **Integration Testing**

Interactions among the modules (data ingestion, models, enforcement pipeline) will be tested to ensure smooth, reliable data flow and communication.

3. System Testing

The complete framework as a whole will be tested to confirm functional and non-functional requirements such as performance, availability, scalability.

4. User Acceptance Testing

The system will be compared with initial project requirements by ensuring that enforcement actions (warnings, restrictions, bans) and justifications are delivered clearly and transparently to the end user.

5 COMMERCIALIZATION AND BUDGETING

- **Commercialization:**
  Can be integrated into existing social media platforms or niche community ecosystems.

- **Investment / Cost:**
  Workforce (developing and testing)
  Cloud storage on Google Cloud (100GB for six months) costing LKR 7200)

- **Recovery:**
  Through a subscription or licensing model for platforms.

- **Target Market:**
  Social media networks, academic communities, and professional networks.

- **Value Proposition:**
  Provides a safer, transparent, more trustworthy and engaging social media environment with explainable AI and emotionally intelligent AI.

- **IPR:**
  Framework for explainable & fair social recommendations can be patented.

REFERENCES

[1] Kepios, "DataReportal," 2025. [Online]. Available: https://datareportal.com/social-media-users. [Accessed 14 August 2025].

[2] Meta, "Fake Accounts," [Online]. Available: https://transparency.meta.com/reports/community-standards-enforcement/fake-accounts/facebook/. [Accessed 15 August 2025].

[3] A. C. G. A. M. P. M. A. S. S. W. K. L. a. K. H. Ann John, "Self-Harm, Suicidal Behaviours, and Cyberbullying in Children and Young People: Systematic Review," *Journal of Medical Internet Research,* vol. 20, no. 4, p. e129, 2018.

[4] Wikipedia, "Social media and suicide," November 2020. [Online]. Available: https://en.wikipedia.org/wiki/Social_media_and_suicide. [Accessed 15 August 2025].

[5] UNICEF, "Cyberbullying: What is it and how to stop it," January 2025. [Online]. Available: https://www.unicef.org/stories/how-to-stop-cyberbullying. [Accessed 15 August 2025].

[6] S. L. CERT|CC, "Sri Lanka Computer Emergency Readiness Team," [Online]. Available: https://www.cert.gov.lk. [Accessed 15 August 2025].

[7] O. f. N. S. (ONS), "Crime in England and Wales: year ending September 2024," Office for National Statistics (ONS), Newport, Wales, United Kingdom, 23 January 2025.

[8] R. I. a. R. DiResta, "How Coordinated Inauthentic Behavior continues on Social Platforms," 29 May 2024. [Online]. Available: https://cyber.fsi.stanford.edu/io/news/how-coordinated-inauthentic-behavior-continues-social-platforms. [Accessed 15 August 2025].

[9] V. o. I. Guy Rosen, "Integrity and Transparency Reports, Third Quarter 2022," 22 November 2022. [Online]. Available: https://about.fb.com/news/2022/11/integrity-and-transparency-reports-q3-2022/.

[10] E. a. V. O. a. D. C. a. M. F. a. F. A. Ferrara, "The rise of social bots," *Communications of the ACM,* vol. 59, no. July 2016, p. 96–104, 2016.

[11] H. a. Z. J. a. H. X. a. H. Z. a. L. A. a. Y. Z. a. Y. P. S. Peng, "Unsupervised Social Bot Detection via Structural Information Theory," *ACM Transactions on Information Systems,* vol. 42, no. 6, pp. 1-42, 2024.

[12] O. V. Yasser Zouzou, "Unsupervised detection of coordinated fake-follower campaigns on social media," *EPJ Data Science,* vol. 13, no. 1, p. 62, 2024.

[13] Y. M. N. Z. G. Z. Y. Q. L. D. S. Q. W. a. X. L. Dong Yuan, "Detecting Fake Accounts in Online Social Networks at the Time of Registrations," in *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security (CCS '19)*, London, UK, 2019.

[14] S. C. A. M. A. V. M. T. Lorenzo Mannocci, "MulBot: Unsupervised Bot Detection Based on Multivariate Time Series," arXiv.org, Cornell University Library, Ithaca, NY, USA, 2022.

[15] A. A. Q. A. A.-H. A. A. Eman Alnagi, "Unmasking Fake Social Network Accounts with Explainable Intelligence," *International Journal of Advanced Computer Science and Applications (IJACSA),* vol. 15, no. 3, pp. 1076-1083, 2024.