An Ethical and Emotionally Intelligent Social Media Plugin Using Responsible & Explainable AI

Abstract—Social media platforms encounter ongoing difficulties with fragmented moderation and interaction operations that do not collaborate effectively to deal with destructive content, inauthentic behavior, emotion-insensitive interaction, and opaque recommendation. This paper features an ethical and emotionally intelligent social media plug-in using responsible and explainable artificial intelligence implemented as a unified real-time service, to social media applications, proven through implementation on the Open-Source Social Network (OSSN). The framework incorporates four elements, namely multimodal cyberbullying detection, behavior-based social bot detection, explainable friend recommendation, and an emotion-aware reaction system. The cyberbullying module is a combination of transformer-based text analysis, image processing, OCR, and keyword fusion to moderate content in an explainable manner, and the fake account module uses temporal behavioral features to detect bot accounts regardless of content. The recommendation system offers interpretable recommendations with context and the emotion-aware module with empathetic interaction via emotion recognition and filtering. Experimental results indicate high performance, of 88.60% on text-based cyberbullying, 68.81% on image moderation, MCC 0.9704 and ROC-AUC 0.9981 on bot detection, and 73.10% F1 on sarcasm detection. The outcomes of deployments demonstrate the potential to have a single, transparent, and human-centric AI system to have safer and more meaningful social media interactions.

Keywords—cyberbullying detection, Explainable artificial intelligence, content moderation, responsible AI, bot detection.

I. INTRODUCTION

Social media platforms currently connect over 5.2 billion users worldwide and are fundamentally reshaping the way individuals communicate, consume information, and how they form communities [1]. Despite this reach, governance of platforms has not kept the pace. The existing paradigms of moderation are largely reactive and fragmented, permitting harmful content, inauthentic accounts, and emotionally intensive interactions to flourish before intervention is possible [2].

Cyberbullying which is increasingly propagating through image-based formats such as memes, causes devastating psychological impact; a 2023 nationally representative survey of the U.S. found that 55% of adolescents had been victims of online victimization [3]. At the same time, automated bot accounts falsify public discourse by using coordinated narratives and their prevalence and influence keep constantly increasing [2]. Beyond explicit harm, platforms also permit emotionally intensive interactions wherein they allow unrestricted reactions to the content irrespective of emotional context. Sarcasm and implicit sentiment further slip past the detection because of their context-dependent nature [4]. Compounding these issues, friend recommendation systems remain opaque, graph-based engines that exclude new users and do not provide any guidance for initiating meaningful interaction [5]. Altogether these challenges erode safety, authenticity, and social quality of online environments, making a call for a unified, proactive, and explainable AI-driven solution.

Although significant advancements have been made in such domains, the existing solutions remain fragmented. Cyberbullying detectors based on transformers have achieved strong text-detecting performance [6], [7] but are blind to image-embedded abuse in memes [8]. Techniques used for bot detection have progressed significantly [9], [10], yet the LLM-powered bots that generate human-like text makes the content-dependent approaches fundamentally fragile [11]. Emotion and Sarcasm models that are benchmarked on TweetEval [12] and GoEmotions [13] focus only on classification and not adaptive to user interactions or provide real-time ethical reaction guidance. Explainable friend recommendations have enhanced the cold start overseeing [5], yet no system combines semantic transparency with proactive interaction guidance [14]. However, there is no current system that integrates safety, authenticity, emotional intelligence, and user engagement as a single deployable and explainable AI framework.

In response to these challenges, this paper proposes a single, proactive AI plug-in that considers all four dimensions by a single integrated service, a combination of special modules in a shared Python inference backend, along with a PHP plugin layer. It is demonstrated through complete real-time deployment on the Open-source Social Network (OSSN) where it operates as a live moderation and engagement service, which validates its real-world applicability beyond benchmark assessment.

The main contributions of this work are:

1. A dual-pipeline multimodal cyberbullying detection that combines transformer-based text analysis and vision-language models for image and meme content, with explainable Allow/Warn/Delete moderation decisions deployed in real time.
2. A content-free, lifecycle-aware social bot detection framework, which uses temporal behavioral features with hybrid machine learning and explainable decision outputs, with demonstrated robustness to the content generated by LLMs.
3. An explainable friend recommendation that provides semantic interest profiling and tone-conditioned interactive guidance which unifies recommendation, per-category explainability, and proactive conversation scaffold together in a single pipeline.
4. A proactive emotion-aware reaction module, which detects emotions, sarcasm and slang across text and image inputs using fine-tuned transformer models, filters available reactions dynamically and provides guidance on empathetic interaction.
5. A unified AI plugin architecture that integrates all four capabilities into a single deployable service, demonstrated through real-time deployment on a live social network.

The rest of this paper is structured as follows: A related work review is in Section II. Section III covers the plugin architecture. Experimental results and evaluation are in section IV. The paper concludes in Section V by summarizing contributions and suggesting recommendations on further research directions.

II. Related work

The modern social media platforms are utilizing artificial intelligence for content moderation, analyze user behavior and improve their interaction. Systems utilized by Meta Platforms (Facebook) involve the combination of natural language processing and computer vision as a solution for moderation of content, yet have limitations related to explainability, contextual interpretation, and dealing with implicit or malicious conduct [7], [15].

Cyberbullying detection has advanced to transformer-based models that develop contextual semantics. Transformer models such as RoBERTa and large language models perform better on learning contextual representations and can detect implicit harmful content [6], [7]. However, multimodal content does not work with text-alone strategies. MemesViTa and CLIP are the techniques that use both textual and visual data to enhance harmful memes and image-based abuse detection [8], [16]. Furthermore, Context-aware models are more accurate but require scaling [4].

Social bot detection has advanced from features and behavioral patterns to deep learning and generalization-based methods [2], [17]. Later, deep learning models incorporated textual and metadata features which enhanced performance but were susceptible to human-like generated texts [18]. Datasets such as TwiBot-20 and TwiBot-22 demonstrate the weaknesses of extrapolation to unknown bot behavior [19], [20], with the modern contributions focusing more on behavior-principled detectors than content-principled ones [11].

The friend recommendation systems have developed beyond the graph-based techniques to semantic and explainable models. Conventional approaches are affected by cold-start problem [5]. Recent literature shows that SBERT embeddings have been used to embed semantic relationships between user interests [21], and explainable recommendation methods enhance user engagement and trust [14]. There has also been an investigation of proactive conversational systems to aid in interaction guidance [22].

Transformer-based emotion detection has improved emotion-sensitive interaction systems. The datasets such as GoEmotions or TweetEval, exhibit good results in extracting fine-grained emotional expressions [12], [13]. There are also multimodal models like CLIP which add more emotion recognition capabilities [16], or explainable AI techniques like LIME which enhance transparency, but only limited to a real-time implementation [23].

Overall, current research reveals growth in specific fields. However, they do not provide a single, coherent system. The proposed system fills this gap with the integration of cyberbullying detection, behavior-based bot detection, explainable recommendation, and emotion-conscious interaction into one ethical AI-friendly social media plug-in.

III. Methodology

A. Proposed Methodology

The proposed system is designed based on certain principles of ethical and responsible artificial intelligence, ensuring fairness, transparency, and accountability, as well as the well-being of the user, as shown in Fig. 1. Fairness is incorporated into the system through proper preprocessing of the data and a balanced design of the model to avoid bias for different types of users. Transparency is also incorporated into the system through the addition of explainable mechanisms, ensuring that the output is interpretable for all the modules used in the system, thereby enhancing user trust, and understanding of the decisions made by the system [23]. Another aspect incorporated into the proposed system is the well-being of the user, ensuring that harmful content, unsafe queries, and inappropriate user interactions are proactively addressed, given the rising issue of platform responsibility and user well-being [1], [15].

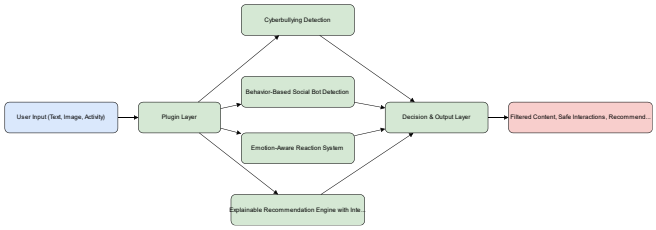

Fig. 1. Overall architecture of the proposed plugin-based social media intelligence system.

B. Data Collection and Preprocessing

The proposed system leverages heterogeneous datasets that incorporate textual data, image data, and behavioral activity logs to ensure that all system components incorporate multimodal and behavioral analysis.

For the detection of cyberbullying, various publicly available datasets such as the Jigsaw Toxic Comments dataset, TRAC dataset, and cyberbullying dataset from the Kaggle platform are integrated into a unified dataset (95 450 samples). The datasets are standardized to ensure that all datasets use unified labels such as mild bullying, bullying, and non-bullying. In addition, datasets such as the MultiOFF dataset and the Hateful Memes dataset are integrated into a combined dataset (9740 samples) for the detection of abusive content using multimodal analysis, which has been proven to improve the detection of harmful content in specific contexts [8].

The preprocessing of textual data includes several steps that ensure the data is clean and normalized into a structured format that can be leveraged by transformer-based models that require a unified input structure for efficient learning.

For image data, resizing and normalization are done to make the data compatible with deep learning architectures. Optical Character Recognition (OCR) is also employed to extract hidden text from the images. This way, the hidden text in memes and images can be analyzed along with the text in the memes. This is important as most harmful intentions in social media are hidden in the context of the image and text combined.

For behavior-based detection, the Cresci-2017 dataset is utilized. The dataset contains five subsets of accounts, ranging from genuine users to bot families. The user activity log data are preprocessed to obtain behavioral features. The behavioral features extracted are posting frequency, time between two consecutive events, activity duration, and burst behavior. These are the most employed features in bot detection research as they detect anomalies in human behavior patterns [2], [9].

For the recommendation component, the user profiles containing interests, hashtags, and text are aggregated. These are combined to form a unified representation. The interests, hashtags, and text are combined to form a single entity. This way, the generation of an effective embedding can be done by employing Sentence-BERT to capture the context between user interests [21].

The preprocessing ensures that all data are converted to a structured and meaningful form, as illustrated in Fig. 2.

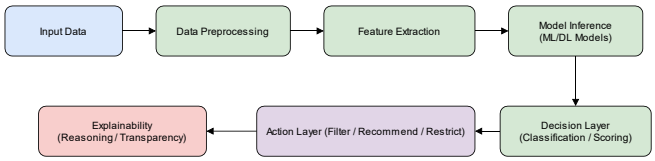

Fig. 2. Common processing pipeline followed by all components in the proposed system.

C. Cyberbullying Detection

The Cyberbullying Detection component uses a multimodal hybrid framework for detecting harmful or abusive content in interactions in social media by categorizing both texts and images. As illustrated in Table I, DistilBERT is utilized in the text pipeline to create contextual embeddings and extract semantic relations in social media language [6], [7]. To enhance the robustness, a keyword-based module with fuzzy matching is employed to identify explicit abusive words and generate an auxiliary score.

In the case of image input, EasyOCR extracts the text appearing as part of the image, and the result is integrated with the user caption in one textual message. CLIP feature extractor is applied to produce joint visual-textual embeddings, and the likelihood of bullying is estimated by means of a Logistic Regression classifier.

The results of the text model, the image model, and the key word model are jointly used through a weighted fusion approach, 0.45 of text, 0.35 of image, and 0.20 of key words. According to the final score, the system assigns one of the three moderation options which are: Allow to post, Warn the user or auto delete the post. An explainability layer enhances transparency by highlighting harmful words, confidence rating, as well as other important contributing factors. As a result, it enhances user trust and helps promote more explainable moderation.

TABLE I. MODELS AND TECHNIQUES USED IN CYBERBULLYING DETECTION.

Technique	Model	Purpose
Text Analysis	DistilBERT	Bullying classification
Image Analysis	EfficientNet-B0	Visual feature extraction
OCR	EasyOCR	Text extraction from images
Fusion	Weighted Score	Multimodal decision making

D. Behavior-Based Social Bot Detection

This framework is a multi-layered anomaly detection approach that analyzes temporal behavioral patterns to identify suspicious or automated accounts, remaining inherently robust against LLM-powered bots that generate human-like textual content.

As summarized in Table II, the first layer processes user activity data to extract seven content-free temporal behavioral features such as account age, activity duration, activity rate, mean and standard deviation of inter-event intervals, coefficient of variation and burst patterns, capturing both the intensity and irregularity of user behavior as effective indicators of automated activity [2], [9].

In the second layer, XGBoost is trained on labeled behavioral data under a strict cross-family generalization protocol. Class imbalance is addressed through weighted training, and the model stability is confirmed by performing five-fold stratified cross-validation.

In the third layer, Isolation Forest models the behavioral distribution of genuine users. A hybrid fusion approach combines supervised and unsupervised outputs through both fixed-weight and confidence-based dynamic weighting, assigning more influence on anomaly detector when the classifier confidence is low, which improves the robustness against novel bot behaviors.

In the fourth layer, early-stage detection limits feature extraction to the first N tweets where optimal N is selected over ablation. A lifecycle-aware layer then integrates early and mature-stage signals dynamically using a normalized Evidence Score based on account age, activity date, and duration.

The final decision layer computes a risk score per account and implements threshold-based moderation measures in the OSSN development environment and Explainability is incorporated through SHapley Additive exPlanations (SHAP) analysis and IQR-based score of behavioral deviation.

TABLE II. LAYERED TECHNIQUES USED IN BEHAVIOR-BASED SOCIAL BOT DETECTION.

Layer	Technique	Purpose
Feature Engineering	Temporal Activity Metrics	Behavioral representation
Supervised Classification	XGBoost (class-weighted)	Baseline bot detection
Anomaly Fusion	Isolation Forest + dynamic α	Novel bot robustness
Early Detection	N-tweet XGBoost + ablation	Pre-history flagging
Lifecycle Integration	Evidence score fusion	Stage-adaptive risk

E. Emotion-Aware Reaction System

The Emotion-Aware Reaction System is proposed as a multimodal, context-aware system that incorporates emotion detection, language understanding, ethical filtering, and user interaction to enable responsible and empathetic user engagement. As summarized in Table III, for text data, a

transformer-based classification approach from the RoBERTa family is applied to classify emotions based on a fine-grained emotion classification scheme with 28 emotion categories, enabling deep emotional understanding of user engagement. For image data, facial emotion recognition is performed by a ResNet50-based model trained on emotion classes like those in FER2013.

For multimodal data, a weighted fusion of text and image models is applied to enable better understanding of user engagement, where text is considered more important, but image information is also utilized to improve accuracy. To improve contextual understanding in a real-world social media environment, a sarcasm detection module and a slang understanding module are integrated into the system.

A rule-based mapping mechanism is proposed to map the identified emotions to a set of allowed emoji reactions, ensuring that user engagement is contextually appropriate and sensitive to other users. In addition, an ethical filtering mechanism is proposed to detect harmful content and unsafe user queries, ensuring user engagement is responsible and empathetic. Moreover, the system provides user interaction guidance by suggesting empathetic responses based on emotional context, transforming the system from a passive classification system to an active decision-support system for socially aware user engagement. Finally, to improve user trust, an explainability feature is incorporated to provide insights into emotion prediction, reaction restriction, and filtering decisions.

F. *Explainable Recommendation Engine with Interaction Guidance*

The Explainable Recommendation Engine, also detailed in Table III, is specifically built to provide recommendations for friend suggestions based on semantic similarity while ensuring transparency and facilitating meaningful user interactions.

User profiles are built, considering the content users actively engage with most frequently by using two approaches: TF-IDF for keyword similarity and Sentence BERT (SBERT) for deeper semantic similarity between the user's interests [21]. The former technique provides a better representation for context and similarity between different topics. The cosine similarity function is used to measure user connection similarity by identifying the top N for most relevant users based on the similarity score. For transparency, the system highlights the common interests between the users and provides human-readable explanations. Moreover, the similarity score for each category is also calculated to show the similarity between the different interests of the users.

The interaction guidance module provides conversation starters based on common interests and the chosen tone for the conversation such as friendly, professional, casual, and icebreaker by applying predefined templates.

This approach extends the traditional recommendation system by not only providing recommendations for connections but also facilitating meaningful interactions with the recommended users, which is in line with the recent advancements in proactive dialogue systems [22].

TABLE III. CORE TECHNIQUES USED IN EMOTION-AWARE REACTION AND RECOMMENDATION SYSTEMS.

Component	Model / Technique	Purpose
Emotion Detection	RoBERTa	Fine-grained emotion classification
Image Emotion	ResNet50	Facial emotion recognition
Reaction Mapping	Rule-based	Emotion-to-emoji control
Recommendation	SBERT	Semantic user similarity

G. *Ethical and Responsible AI*

The proposed system has been designed based on a few important principles of ethical and responsible AI, including fairness, transparency, accountability, and user well-being. Fairness in the proposed system can be achieved through proper preprocessing of the data. In addition, a balanced model can be used to minimize the level of bias between different types of users. Transparency in the proposed system can be achieved through the addition of features to the system to ensure transparency in the decision process. This will help to increase the level of trust between the system and the user. In addition, accountability in the proposed system can be achieved through transparency in the decision process. Furthermore, user well-being in the proposed system can be achieved through proper identification of harmful content, unsafe queries, and user interaction, which will help to solve the user safety issue, as mentioned in [1], [15].

IV. RESULTS AND EVALUATION

Each capability was evaluated independently on task-specific benchmarks, followed by unified deployment evaluation on OSSN. Primary measures for classification tasks are Accuracy and macro F1, whereas Mathews Correlation Coefficient (MCC) and ROC-AUC are used for bot detection due to class imbalance.

A. *Cyberbullying Detection System*

Text and image pipelines were trained on multiple datasets with a 70/15/15 stratified split, were evaluated on the test datasets. Models including ResNet and MobileNet-V2 discarded on account of poor generalization performance (val. acc. 54.80%). The weighted-average test results are presented in Table IV. DistilBERT showed the highest results in text classification with accuracy of 88.30 and F1-score of 88.29, unlike BERT (110M) and RoBERTa (125M), demonstrating high accuracy and significantly reduced memory consumption and inference time, which allowed it to be deployed in real-time. In image-based detection, CLIP using Logistic Regression performed surpassed EfficientNet-B0 (68.81% vs. 51.40) through combined visual-textual representations. Despite average performance, it is well-integrated in a weighted fusion system (text: 0.45, image: 0.35, keyword: 0.20) combined with a severe-word detection system, facilitating immediate elimination of content with high risks.

TABLE IV. MODEL EVALUATION METRIS

Model	Accuracy	Precision	Recall	F1	selected
DistilBERT (text)	88.62%	88.29%	88.30%	88.29%	✓
BERT (text)	88.10%	88.35%	88.10%	88.19%	✗

ELECTRA (text)	86.28%	87.93%	86.28%	86.58%	✗
RoBERTa (text)	87.07%	88.11%	87.07%	87.30%	✗
CLIP + LR (image)	68.81%	70.00%	69.00%	69.00%	✓
EfficientNet-B0 (image)	51.40%	51.40%	51.40%	51.39%	✗

B. Behaviour-Based Social Bot Detection

The framework was evaluated in a cross-family generalization protocol. Models trained on three seen Cresci-2017 subsets and were tested on two entirely unseen subsets. (Spambots 2, Spambots 3). Table V summarizes the cross-family generalization results across all layers. XGBoost (L2) achieves MCC 0.9769 on Spambots 2 with no post content, profiles, or graph structure. The fixed weight fusion (MCC 0.6432) is outperformed by L3 dynamic alpha fusion (MCC 0.6995) on Spambots 3. A moderate performance on Spambots 3 (MCC 0.4819) is indicated by L4 early detection (N=20, selected by ablation study) which motivates L5 lifecycle fusion, achieving MCC 0.9704 and ROC-AUC 0.9981. This reflects a deliberate early-stage trade off. The low L5 MCC on Spambots 3 (0.5635) shows that it is a deliberate mimicry of human temporal behavior and ROC-AUC 0.9596 confirms its strong class separability.

TABLE V. CROSS-FAMILY GENERALIZATION ON CRESCI-2017 DATASET

Layer / Model	Evaluation Set	MCC	ROC-AUC	F1
L2 - XGBoost (class-weighted)	Spambots 2 (unseen)	0.9769	0.9980	0.99
L2 - XGBoost (class-weighted)	Spambots 3 (unseen)	0.7802	0.9682	0.88
L3 - Hybrid Dynamic alpha	Spambots 2 (unseen)	0.9736	0.9985	0.99
L3 - Hybrid Dynamic alpha	Spambots 3 (unseen)	0.6995	0.9650	0.83
L4 - Early Detection (N=20)	Spambots 2 (unseen)	0.6735	0.8590	0.84
L4 - Early Detection (N=20)	Spambots 3 (unseen)	0.4819	0.7263	0.67
L5 - Lifecycle Fusion (deployed)	Spambots 2 (unseen)	0.9704	0.9981	0.99
L5 - Lifecycle Fusion (deployed)	Spambots 3 (unseen)	0.5635	0.9596	0.72

C. Emotion-Aware Reaction System

Three sub-tasks were assessed separately, i.e., facial emotion recognition (AffectNet, 8-class), sarcasm detection, and fine-grained text emotion classification (GoEmotions, 28-label), and optimal models are cited in Table VI. EfficientNet-B2 has moderate Top-1 (62.64%) and strong Top-2 (78.66%) accuracy, Macro AUC (0.895), and Micro AUC (0.903), attesting to the fact that near-miss predictions are semantically proximate to ground truth (good enough to achieve coarse valence detection in the reaction-filtering pipeline). Fine-tuned RoBERTa has F1 73.10% in sarcasm detection. With a 28-label classifier, the accuracy of 46.87% indicates the challenge of detecting emotions with fine-grain. The greater accuracy (63.30%) compared to recall (41.94) is appropriate in the reaction filtering, where false positives are of utmost concern.

TABLE VI. TASK SPECIFC RESULTS

Task	Model	Primary Metric	Value	Secondary Metrics	Value
Facial Emotion (8-class)	Efficient Net-B2	Top-1 Acc.	62.64 %	Top-2 / Macro AUC / Micro AUC	78.66% / 0.895 / 0.903
Sarcasm Detection	Fine-tuned RoBERTa	F1	73.10 %	Accuracy / Eval Loss	73.09% / 0.542
Text Emotion (28-label)	RoBERTa-base	Accuracy	46.87 %	F1 / Prec. / Rec.	46.94% / 63.30% / 41.94%

D. Explainable Friend Recommendation with Interaction Guidance

Evaluated on a social media dataset of 20,000 users where candidates are ranked with SBERT cosine similarity. SBERT captures semantic relations between differently expressed interests such as 'AI' and 'machine learning' and pairs that scored above 0.65 are retained as recommendations per the qualification schema in Table VII. Similarity scores by category (Music: 80%, Art: 90%) are intrinsically derived from recommendation mechanism and not post-hoc, satisfying XAI faithful requirements. Cold-start problem is resolved by profile-only encoding for any user with at least one declared interest. A proof-of-concept user study (n=20, five-point Likert scale) clarified that 80-85% of participants agreed explanations on recommendation improved understanding, and that 75-80% found tone conditioned conversation starters (64 templates: 16 categories x 4 tones) were useful in starting interaction.

TABLE VII. SIMILARITY SCORE QUALIFICATION SCHEMA

Similarity Score Range	Recommendation Quality
< 0.50	Weak - dissimilar interest profiles
0.50 – 0.64	Moderate - partial interest overlap
0.65 – 0.70	High - semantically relevant matches
> 0.70	Very high - strongly aligned profiles

E. Comparative Analysis

Table VIII benchmarks the proposed plugin against prior work across seven dimensions. No current system addresses all four capabilities simultaneously with integrated XAI and live deployment. This plugin unifies multimodal cyberbullying detection, content-free bot detection robust to LLM-generated text, explainable friend recommendation with tone-guided interactions, and proactive emotion-aware reaction control in a single deployable framework in real-time.

TABLE VIII. CAPABILITY COVERAGE COMPARISON WITH PRIOR WORK

	[6]	[8]	[10]	[5]	[12]	Proposed Plugin
Cyberbullying Detection	✓	✓	✗	✗	✗	✓
Bot Account Detection	✗	✗	✓	✗	✗	✓
Friend Recommendation	✗	✗	✗	✓	✗	✓
Emotion & Sarcasm Detection	✗	✗	✗	✗	✓	✓
Explainability (XAI)	✗	✗	✗	✓	✗	✓
Live Deployment	✗	✗	✗	✗	✗	✓
Unified Component Framework	✗	✗	✗	✗	✗	✓

V. Conclusion and Future Work

This paper presented a multi-component ethical and explainable AI plugin for the OSSN platform, integrating Cyberbullying Detection, Behavior-Based Social Bot Detection, Emotion-Aware Reaction System, and Explainable Friend Recommendation System with Interaction Guidance. Cyberbullying detection achieved 88.62% accuracy on text-based content with a hybrid multimodal approach, extending the detection to meme-based abuse, while bot detection showed a strong cross-family generalization and achieved (MCC 0.9704, ROC-AUC 0.9981) using content-free temporal features, demonstrating robustness against unseen automated accounts. The emotion-aware reaction system was able to limit emotionally inappropriate engagements using transformer-based emotion and sarcasm recognition, and the friend suggestion engine offered semantically similar suggestions with per-class clarification and similarity-based, tone conscious connection direction. The process of explainability to all the four elements makes sure that users and administrators can comprehend and trust the decisions made by the system, which is one of the most important gaps in the existing social media moderation strategies.

Future work should focus on validating all components on more recent and larger datasets to improve generalization of modern social media behavior. Additionally, optimizing components for real-time processing would improve responsiveness. As Explainability layers currently operate in varying depths across different components, future work should explore unified explainability standards across the plugin. Furthermore, capability to video and audio detection would result in a system that identify more harmful content.

Reference

[1] S. Kemp, "Digital 2024: Global Overview Report," DataReportal, We Are Social, Meltwater, Jan. 2024. [Online]. Available: https://datareportal.com/reports/digital-2024-global-overview-report. [Accessed: Mar. 22, 2026].

[2] O. Varol, E. Ferrara, C. Davis, F. Menczer, and A. Flammini, "Online human-bot interactions: Detection, estimation, and characterization," in Proc. Int. AAAI Conf. Web Soc. Media, 2017.

[3] J. W. Patchin and S. Hinduja, "2023 Cyberbullying Data," Cyberbullying Research Center, Feb. 2024. [Online]. Available: https://cyberbullying.org/2023-cyberbullying-data. [Accessed: Mar. 22, 2026].

[4] J. M. Pérez, F. M. Luque, D. Zayat, M. Kondratzky, and A. Moro, "Assessing the impact of contextual information in hate speech detection," IEEE Access, vol. 11, pp. 30575–30590, 2023.

[5] C. Li, B. Zhou, W. Lin, Z. Tang, Y. Tang, Y. Yong, and J. Cao, "A personalized explainable learner implicit friend recommendation method," Data Sci. Eng., vol. 8, no. 1, pp. 23–35, 2023.

[6] A. A. Jamjoom, H. Karamti, M. Umer, S. Alsubai, T.-H. Kim, and I. Ashraf, "Robertanet: Enhanced RoBERTa transformer based model for cyberbullying detection with GloVe features," IEEE Access, vol. 12, pp. 58950–58959, 2024.

[7] A. Albladi et al., "Hate speech detection using large language models: A comprehensive review," IEEE Access, vol. 13, pp. 20871–20892, 2025.

[8] T. Sultan, M. A. T. Rony, M. S. Islam, S. Aldosary, and W. El-Shafai, "Memesvita: A novel multimodal fusion technique for troll memes identification," IEEE Access, vol. 12, pp. 177811–177828, 2024.

[9] M. Aljabri et al., "Machine learning-based social media bot detection: A comprehensive literature review," Soc. Netw. Anal. Mining, vol. 13, no. 1, p. 20, 2023.

[10] K. Hayawi et al., "Social media bot detection with deep learning methods: A systematic review," Neural Comput. Appl., vol. 35, no. 12, pp. 8903–8918, 2023.

[11] E. Ferrara, "Social bot detection in the age of ChatGPT: Challenges and opportunities," First Monday, 2023.

[12] F. Barbieri, J. Camacho-Collados, L. E. Anke, and L. Neves, "TweetEval: Unified benchmark and comparative evaluation for tweet classification," in Findings Assoc. Comput. Linguistics: EMNLP, 2020.

[13] D. Demszky et al., "GoEmotions: A dataset of fine-grained emotions," in Proc. 58th Annu. Meeting Assoc. Comput. Linguistics, 2020.

[14] J. Govea, R. Gutierrez, and W. Villegas-Ch, "Transparency and precision in the age of AI: Evaluation of explainability-enhanced recommendation systems," Front. Artif. Intell., vol. 7, p. 1410790, 2024.

[15] Meta Platforms, "Community Standards Enforcement Report," 2023. [Online]. Available: https://transparency.fb.com/reports/community-standards-enforcement/. [Accessed: Mar. 22, 2026].

[16] A. Radford et al., "Learning transferable visual models from natural language supervision," in Proc. 38th Int. Conf. Mach. Learn. (ICML), vol. 139, pp. 8748–8763, 2021.

[17] S. Cresci et al., "The paradigm-shift of social spambots: Evidence, theories, and tools for the arms race," in Proc. 26th Int. World Wide Web Companion Conf. (WWW Companion), pp. 963–972, 2017.

[18] C. Kudugunta and E. Ferrara, "Deep neural networks for bot detection," Inf. Sci., vol. 467, pp. 312–322, 2018.

[19] S. Feng et al., "TwiBot-20: A comprehensive Twitter bot detection benchmark," in Proc. 30th ACM Int. Conf. Inf. Knowl. Manage. (CIKM), pp. 4485–4494, 2021.

[20] S. Feng et al., "TwiBot-22: Towards graph-based Twitter bot detection," Adv. Neural Inf. Process. Syst., vol. 35, pp. 35254–35269, 2022.

[21] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in Proc. EMNLP-IJCNLP, pp. 3982–3992, 2019.

[22] Y. Deng, W. Lei, W. Lam, and T.-S. Chua, "A survey on proactive dialogue systems: Problems, methods, and prospects," arXiv preprint arXiv:2209.12602, 2022.

[23] M. T. Ribeiro, S. Singh, and C. Guestrin, "Why should I trust you? Explaining the predictions of any classifier," in Proc. 22nd ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining (KDD), pp. 1135–1144, 2026